

ista_solve_hot

April 18, 2022

```
[54]: import numpy as np
from scipy.io import loadmat
import matplotlib.pyplot as plt
from sklearn.svm import LinearSVC
import numpy.random
import math
```

```
[151]: def ista_solve_hot( A, d, la_array ):
    # ista_solve_hot: Iterative soft-thresholding for multiple values of
    # lambda with hot start for each case - the converged value for the previous
    # value of lambda is used as an initial condition for the current lambda.
    # this function solves the minimization problem
    # Minimize  $\|Ax-d\|_2^2 + \lambda\|x\|_1$  (Lasso regression)
    # using iterative soft-thresholding.
    max_iter = 10**4
    tol = 10**(-3)
    tau = 1/np.linalg.norm(A,2)**2
    n = A.shape[1]
    w = np.zeros((n,1))
    num_lam = len(la_array)
    X = np.zeros((n, num_lam))
    for i, each_lambda in enumerate(la_array):
        for j in range(max_iter):
            z = w - tau*(A.T@(A@w-d))
            w_old = w
            w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
            X[:, i:i+1] = w
            if np.linalg.norm(w - w_old) < tol:
                break
    return X
```

```
[242]: in_data = loadmat('BreastCancer.mat')
X = in_data['X']
y = in_data['y']
```

```
[243]: print(X.shape, y.shape)
```

(295, 8141) (295, 1)


a)

```
[254]: # We will use the first 100 patients
X_train = X[:100]
y_train = y[:100]
# Set lambda array with range (1e-6, 20) spaced logarithmically
lambda_array = []
i = 1e-6
while i < 19:
    lambda_array.append(i)
    i += math.log(2)
lambda_array.append(20)
w = ista_solve_hot(X_train, y_train, lambda_array)
```

```
[245]: residual_error = []
w_norm = []
for i in range(len(lambda_array)):
    residual_error.append(np.linalg.norm(X_train@w[:,i] - y_train))
    w_norm.append(abs(w[:,i]).sum())
    print("w_norm: " , w_norm[i], "\tresidual_error: ", residual_error[i], "\tlambda: ", lambda_array[i])
plt.scatter(x=w_norm, y=residual_error)
plt.show()
# As lambda increases the residual error and the norm of w decreases
```


| | | |
|----------------------------|------------------------------------|----------------------------|
| w_norm: 36.109055952630406 | residual_error: 122.62017023893092 | lambda: 1e-06 |
| w_norm: 18.1266980593891 | residual_error: 117.3028693792493 | lambda: 0.6931481805599453 |
| w_norm: 13.51150995902426 | residual_error: 113.75422580631017 | lambda: 1.3862953611198905 |
| w_norm: 11.207473857087024 | residual_error: 110.96552155127395 | lambda: 2.079442541679836 |
| w_norm: 9.600042255143801 | residual_error: 108.64335521508443 | lambda: 2.7725897222397813 |
| w_norm: 8.425137987735571 | residual_error: 106.67098689091421 | lambda: 3.4657369027997267 |
| w_norm: 7.504537201776889 | residual_error: 104.93988316904763 | lambda: 4.158884083359672 |
| w_norm: 6.785326085022339 | residual_error: 103.44476518158844 | lambda: 4.852031263919617 |
| w_norm: 6.123578969469852 | residual_error: 102.12035118812786 | lambda: 5.545178444479562 |
| w_norm: 5.579710963832548 | residual_error: 100.97384444664706 | lambda: 6.238325625039508 |
| w_norm: 5.137386633432994 | residual_error: 99.97513990804859 | lambda: 6.931472805599453 |
| w_norm: 4.7294810562842615 | residual_error: 99.12557162414309 | lambda: 7.624619986159399 |

| | | |
|----------------------------|-----------------------------------|---------|
| w_norm: 4.355999922114707 | residual_error: 98.32365768967908 | lambda: |
| 8.317767166719344 | | |
| w_norm: 3.9723493979938773 | residual_error: 97.50426863953899 | lambda: |
| 9.01091434727929 | | |
| w_norm: 3.75616629710407 | residual_error: 97.0082963354693 | lambda: |
| 9.704061527839235 | | |
| w_norm: 3.4778926816773117 | residual_error: 96.27848306896577 | lambda: |
| 10.39720870839918 | | |
| w_norm: 3.136669482594097 | residual_error: 95.45134190969435 | lambda: |
| 11.090355888959126 | | |
| w_norm: 2.9903158145972366 | residual_error: 95.0538750806501 | lambda: |
| 11.783503069519071 | | |
| w_norm: 2.764883558007083 | residual_error: 94.46313591228582 | lambda: |
| 12.476650250079016 | | |
| w_norm: 2.5863002243014237 | residual_error: 93.99688909430621 | lambda: |
| 13.169797430638962 | | |
| w_norm: 2.4504933690532065 | residual_error: 93.6815802149507 | lambda: |
| 13.862944611198907 | | |
| w_norm: 2.2633392534203525 | residual_error: 93.24543582959832 | lambda: |
| 14.556091791758853 | | |
| w_norm: 2.0540384197405395 | residual_error: 92.78511251484086 | lambda: |
| 15.249238972318798 | | |
| w_norm: 1.956318032917396 | residual_error: 92.6079892263537 | lambda: |
| 15.942386152878743 | | |
| w_norm: 1.8420240652914215 | residual_error: 92.39630102542934 | lambda: |
| 16.635533333438687 | | |
| w_norm: 1.7578344358851319 | residual_error: 92.2403313247562 | lambda: |
| 17.32868051399863 | | |
| w_norm: 1.6912694674229027 | residual_error: 92.11599146101157 | lambda: |
| 18.021827694558574 | | |
| w_norm: 1.5822541740103822 | residual_error: 91.92717363184512 | lambda: |
| 18.714974875118518 | | |
| w_norm: 1.426055638400412 | residual_error: 91.72082256986741 | lambda: |
| 20 | | |



b)

```
[246]: error_rate = []
        sparsity = np.zeros(len(lambda_array))
        for i in range(len(lambda_array)):
            diff = abs(np.sign(X_train@w[:,i]) - y_train[:,0])
            for j in w[:,i]:
                # Add to count the number of nonzero entries of w => if |w|_i > 1e-6
                if j > 1e-6:
                    sparsity[i] += 1
            error_rate.append(diff.sum()/len(diff))
        plt.scatter(x=sparsity, y=error_rate)
        plt.xlabel("Sparsity")
        plt.ylabel("Error rate")
        plt.show()
        # As the number of nonzero entries in w increases => the error rate decreases
```



c)


```
[247]: X_test = X[101:]
       y_test = y[101:]
```

```
[253]: # a)
       residual_error = []
       w_norm = []
       for i in range(len(lambda_array)):
           residual_error.append(np.linalg.norm(X_test@w[:,i] - y_test))
           w_norm.append(abs(w[:,i]).sum())
           print("w_norm: " , w_norm[i], "\tresidual_error: ", residual_error[i],
               ↳ "\tlambda: ", lambda_array[i])
       plt.scatter(x=w_norm, y=residual_error)
       plt.show()
       # As with the train data, as lambda increases the residual error and the norm-1
       ↳ of w decreases.
       # However, the plot show a more curved graph with it's point with a more
       ↳ distance between them
```

| | | |
|----------------------------|------------------------------------|----------------------------|
| w_norm: 36.109055952630406 | residual_error: 201.21304378897253 | lambda: 1e-06 |
| w_norm: 18.1266980593891 | residual_error: 197.91570675411856 | lambda: 0.6931481805599453 |
| w_norm: 13.51150995902426 | residual_error: 195.93637747867436 | lambda: 1.3862953611198905 |


| | | |
|--|------------------------------------|---------|
| w_norm: 11.207473857087024 2.079442541679836 | residual_error: 194.34688419042368 | lambda: |
| w_norm: 9.600042255143801 2.7725897222397813 | residual_error: 192.82475715297167 | lambda: |
| w_norm: 8.425137987735571 3.4657369027997267 | residual_error: 191.78903343197146 | lambda: |
| w_norm: 7.504537201776889 4.158884083359672 | residual_error: 190.782753642679 | lambda: |
| w_norm: 6.785326085022339 4.852031263919617 | residual_error: 189.77979622919315 | lambda: |
| w_norm: 6.123578969469852 5.545178444479562 | residual_error: 188.87130033785922 | lambda: |
| w_norm: 5.579710963832548 6.238325625039508 | residual_error: 187.91498643542022 | lambda: |
| w_norm: 5.137386633432994 6.931472805599453 | residual_error: 187.08197273897952 | lambda: |
| w_norm: 4.7294810562842615 7.624619986159399 | residual_error: 186.45962974159804 | lambda: |
| w_norm: 4.355999922114707 8.317767166719344 | residual_error: 185.78010100452067 | lambda: |
| w_norm: 3.9723493979938773 9.01091434727929 | residual_error: 185.0836551926895 | lambda: |
| w_norm: 3.75616629710407 9.704061527839235 | residual_error: 184.57120572253032 | lambda: |
| w_norm: 3.4778926816773117 10.39720870839918 | residual_error: 183.74603478027777 | lambda: |
| w_norm: 3.136669482594097 11.090355888959126 | residual_error: 182.7934822301498 | lambda: |
| w_norm: 2.9903158145972366 11.783503069519071 | residual_error: 182.30658498752715 | lambda: |
| w_norm: 2.764883558007083 12.476650250079016 | residual_error: 181.60219752866593 | lambda: |
| w_norm: 2.5863002243014237 13.169797430638962 | residual_error: 181.06824398874377 | lambda: |
| w_norm: 2.4504933690532065 13.862944611198907 | residual_error: 180.7230201718312 | lambda: |
| w_norm: 2.2633392534203525 14.556091791758853 | residual_error: 180.17953711686815 | lambda: |
| w_norm: 2.0540384197405395 15.249238972318798 | residual_error: 179.62115975839671 | lambda: |
| w_norm: 1.956318032917396 15.942386152878743 | residual_error: 179.41276239881375 | lambda: |
| w_norm: 1.8420240652914215 16.63553333438687 | residual_error: 179.11612841938015 | lambda: |
| w_norm: 1.7578344358851319 17.32868051399863 | residual_error: 178.87761119660146 | lambda: |
| w_norm: 1.6912694674229027 18.021827694558574 | residual_error: 178.69479609154843 | lambda: |

w_norm: 1.5822541740103822 residual_error: 178.4147295472156 lambda:
 18.714974875118518
 w_norm: 1.426055638400412 residual_error: 178.11348094047474 lambda:
 20



```

[252]: # b)
error_rate = []
sparsity = np.zeros(len(lambda_array))
for i in range(len(lambda_array)):
    diff = abs(np.sign(X_test@w[:,i]) - y_test[:,0])
    for j in w[:,i]:
        # Add to count the number of nonzero entries of w => if |w|_i > 1e-6
        if j > 1e-6:
            sparsity[i] += 1
    error_rate.append(diff.sum()/len(diff))
plt.scatter(x=sparsity, y=error_rate)
plt.xlabel("Sparsity")
plt.ylabel("Error rate")
plt.show()
# Now, the error rate vs sparsity behaves completely different as before. As
# the sparsity increases, the error
# rate increases with it.
  
```



Q2_starter

April 18, 2022

```
[7]: ## Breast Cancer LASSO Exploration
     ## Prepare workspace
     from scipy.io import loadmat
     import numpy as np
     import math
     X = loadmat("BreastCancer.mat")['X']
     y = loadmat("BreastCancer.mat")['y']
```

```
[8]: def ista_solve_hot( A, d, la_array ):
     # ista_solve_hot: Iterative soft-thresholding for multiple values of
     # lambda with hot start for each case - the converged value for the previous
     # value of lambda is used as an initial condition for the current lambda.
     # this function solves the minimization problem
     # Minimize  $\|Ax-d\|_2^2 + \lambda\|x\|_1$  (Lasso regression)
     # using iterative soft-thresholding.
     max_iter = 10**4
     tol = 10**(-3)
     tau = 1/np.linalg.norm(A,2)**2
     n = A.shape[1]
     w = np.zeros((n,1))
     num_lam = len(la_array)
     X = np.zeros((n, num_lam))
     for i, each_lambda in enumerate(la_array):
         for j in range(max_iter):
             z = w - tau*(A.T@(A@w-d))
             w_old = w
             w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
             X[:, i:i+1] = w
             if np.linalg.norm(w - w_old) < tol:
                 break
         return X
     lam_vals = []
     i = 1e-6
     while i < 19:
         lam_vals.append(i)
         i += math.log(2)
     lam_vals.append(20)
```

```

[78]: ## 10-fold CV

# each row of setindices denotes the starting an ending index for one
# partition of the data: 5 sets of 30 samples and 5 sets of 29 samples
setindices = 
    → [[1,30],[31,60],[61,90],[91,120],[121,150],[151,179],[180,208],[209,237],[238,266],[267,295]

# each row of holdoutindices denotes the partitions that are held out from
# the training set
holdoutindices = [[1,2],[2,3],[3,4],[4,5],[5,6],[7,8],[9,10],[10,1]]

cases = len(holdoutindices)

# be sure to initiate the quantities you want to measure before looping
# through the various training, validation, and test partitions
#
#
#
errors = []
W = []

# Loop over various cases
for j in range(cases):
    # row indices of first validation set
    v1_ind = np.
    → arange(setindices[holdoutindices[j][0]-1][0]-1,setindices[holdoutindices[j][0]-1][1])

    # row indices of second validation set
    v2_ind = np.
    → arange(setindices[holdoutindices[j][1]-1][0]-1,setindices[holdoutindices[j][1]-1][1])

    # row indices of training set
    trn_ind = list(set(range(295))-set(v1_ind)-set(v2_ind))

    # define matrix of features and labels corresponding to first
    # validation set
    Av1 = X[v1_ind,:]
    bv1 = y[v1_ind]

    # define matrix of features and labels corresponding to second
    # validation set
    Av2 = X[v2_ind,:]
    bv2 = y[v2_ind]

    # define matrix of features and labels corresponding to the
    # training set
    At = X[trn_ind,:]

```

```

bt = y[trn_ind]

#     print(len(v1_ind), len(v2_ind), len(trn_ind))
w_case = ista_solve_hot(At,bt,lam_vals) # w matrix of w's for each lambda
error_case = []
for i in range(len(lam_vals)):
    error_case.append( (np.sign(Av1@w_case[:,i]) - bv1).sum() /len(bv1))
errors.append(error_case)
W.append(w_case)

# Use training data to learn classifier
# W = ista_solve_hot(At,bt,lam_vals)

```

```

[79]: # Find best lambda value using first validation set, then evaluate
# performance on second validation set, and accumulate performance metrics
# over all cases partitions
# We search for the best lambda
min_case_idx = None
min_error = None
min_lambda = None
for i in range(len(errors)): # iterate through each case => i=0 is case of [1,2]
    change_lambda = False
    for j in range(len(errors[i])):
        if min_error == None or errors[i][j] <= min_error:
            min_error = errors[i][j]
            min_lambda = lam_vals[j]
            change_lambda = True
    if change_lambda:
        min_case_idx = i

```

```

[80]: # Calculate with the best case of holdout and lambda
v1_ind = np.
    → arange(setindices[holdoutindices[i][0]-1][0]-1,setindices[holdoutindices[i][0]-1][1])

v2_ind = np.
    → arange(setindices[holdoutindices[i][1]-1][0]-1,setindices[holdoutindices[i][1]-1][1])

trn_ind = list(set(range(295))-set(v1_ind)-set(v2_ind))

Av1 = X[v1_ind,:]
bv1 = y[v1_ind]

Av2 = X[v2_ind,:]
bv2 = y[v2_ind]

At = X[trn_ind,:]
bt = y[trn_ind]

```

```

w = ista_solve_hot(At,bt,[min_lambda])

y1_pred = np.sign(Av1@w) - bv1
y2_pred = np.sign(Av2@w) - bv2

y_pred = np.concatenate((y1_pred, y2_pred))

misclassifications = 0
for i in y_pred:
    if i != 0:
        misclassifications += 1

error_rate = misclassifications / len(y_pred)
squared_error = (np.linalg.norm(Av1@w - bv1))**2 + (np.linalg.norm(Av2@w -
→bv2))**2

print("Error rate: " , error_rate)
print("Squared error: ", squared_error)

```

Error rate: 0.3898305084745763
 Squared error: 55.86164056908563

```

[99]: length = len(At.transpose())
      lambdas_matrix = np.zeros((length,length))
      np.fill_diagonal(lambdas_matrix, min_lambda)

```

```

[101]: w_ridge = np.linalg.inv(At.transpose()@At + lambdas_matrix)@At.transpose()@bt
      w_ridge

```

```

[101]: array([[ -0.00472367],
              [ 0.00706112],
              [ 0.00881967],
              ...,
              [-0.00279695],
              [-0.00814373],
              [-0.00346693]])

```

```

[105]: y1_pred_ridge = np.sign(Av1@w_ridge) - bv1
      y1_pred_ridge = np.sign(Av2@w_ridge) - bv2

      y_pred_ridge = np.concatenate((y1_pred_ridge, y1_pred_ridge))

      misclassifications_ridge = 0
      for i in y_pred_ridge:
          if i != 0:
              misclassifications_ridge += 1

```

```
error_rate_ridge = misclassifications_ridge / len(y_pred_ridge)
squared_error_ridge = (np.linalg.norm(Av1@w_ridge - bv1))**2 + (np.linalg.
    ↪norm(Av2@w_ridge - bv2))**2

print("Error rate: " , error_rate_ridge)
print("Squared error: ", squared_error_ridge)
```

Error rate: 0.2

Squared error: 62.012381589953506

We can see that the error rate is lower if we use ridge regression but the squared error is higher

[]: