

ista_solve_hot

April 18, 2022

```
[54]: import numpy as np
      from scipy.io import loadmat
      import matplotlib.pyplot as plt
      from sklearn.svm import LinearSVC
      import numpy.random
      import math
```

```
[151]: def ista_solve_hot( A, d, la_array ):
      # ista_solve_hot: Iterative soft-thresholding for multiple values of
      # lambda with hot start for each case - the converged value for the previous
      # value of lambda is used as an initial condition for the current lambda.
      # this function solves the minimization problem
      # Minimize  $\|Ax-d\|_2^2 + \lambda\|x\|_1$  (Lasso regression)
      # using iterative soft-thresholding.
      max_iter = 10**4
      tol = 10**(-3)
      tau = 1/np.linalg.norm(A,2)**2
      n = A.shape[1]
      w = np.zeros((n,1))
      num_lam = len(la_array)
      X = np.zeros((n, num_lam))
      for i, each_lambda in enumerate(la_array):
          for j in range(max_iter):
              z = w - tau*(A.T@(A@w-d))
              w_old = w
              w = np.sign(z) * np.clip(np.abs(z)-tau*each_lambda/2, 0, np.inf)
              X[:, i:i+1] = w
              if np.linalg.norm(w - w_old) < tol:
                  break
      return X
```

```
[242]: in_data = loadmat('BreastCancer.mat')
      X = in_data['X']
      y = in_data['y']
```

```
[243]: print(X.shape, y.shape)
```

(295, 8141) (295, 1)


a)

```
[254]: # We will use the first 100 patients
X_train = X[:100]
y_train = y[:100]
# Set lambda array with range (1e-6, 20) spaced logarithmically
lambda_array = []
i = 1e-6
while i < 19:
    lambda_array.append(i)
    i += math.log(2)
lambda_array.append(20)
w = ista_solve_hot(X_train, y_train, lambda_array)
```

```
[245]: residual_error = []
w_norm = []
for i in range(len(lambda_array)):
    residual_error.append(np.linalg.norm(X_train@w[:,i] - y_train))
    w_norm.append(abs(w[:,i]).sum())
    print("w_norm: " , w_norm[i], "\tresidual_error: ", residual_error[i], "\tlambda: ", lambda_array[i])
plt.scatter(x=w_norm, y=residual_error)
plt.show()
# As lambda increases the residual error and the norm of w decreases
```


w_norm: 36.109055952630406	residual_error: 122.62017023893092	lambda: 1e-06
w_norm: 18.1266980593891	residual_error: 117.3028693792493	lambda: 0.6931481805599453
w_norm: 13.51150995902426	residual_error: 113.75422580631017	lambda: 1.3862953611198905
w_norm: 11.207473857087024	residual_error: 110.96552155127395	lambda: 2.079442541679836
w_norm: 9.600042255143801	residual_error: 108.64335521508443	lambda: 2.7725897222397813
w_norm: 8.425137987735571	residual_error: 106.67098689091421	lambda: 3.4657369027997267
w_norm: 7.504537201776889	residual_error: 104.93988316904763	lambda: 4.158884083359672
w_norm: 6.785326085022339	residual_error: 103.44476518158844	lambda: 4.852031263919617
w_norm: 6.123578969469852	residual_error: 102.12035118812786	lambda: 5.545178444479562
w_norm: 5.579710963832548	residual_error: 100.97384444664706	lambda: 6.238325625039508
w_norm: 5.137386633432994	residual_error: 99.97513990804859	lambda: 6.931472805599453
w_norm: 4.7294810562842615	residual_error: 99.12557162414309	lambda: 7.624619986159399

w_norm: 4.355999922114707	residual_error: 98.32365768967908	lambda:
8.317767166719344		
w_norm: 3.9723493979938773	residual_error: 97.50426863953899	lambda:
9.01091434727929		
w_norm: 3.75616629710407	residual_error: 97.0082963354693	lambda:
9.704061527839235		
w_norm: 3.4778926816773117	residual_error: 96.27848306896577	lambda:
10.39720870839918		
w_norm: 3.136669482594097	residual_error: 95.45134190969435	lambda:
11.090355888959126		
w_norm: 2.9903158145972366	residual_error: 95.0538750806501	lambda:
11.783503069519071		
w_norm: 2.764883558007083	residual_error: 94.46313591228582	lambda:
12.476650250079016		
w_norm: 2.5863002243014237	residual_error: 93.99688909430621	lambda:
13.169797430638962		
w_norm: 2.4504933690532065	residual_error: 93.6815802149507	lambda:
13.862944611198907		
w_norm: 2.2633392534203525	residual_error: 93.24543582959832	lambda:
14.556091791758853		
w_norm: 2.0540384197405395	residual_error: 92.78511251484086	lambda:
15.249238972318798		
w_norm: 1.956318032917396	residual_error: 92.6079892263537	lambda:
15.942386152878743		
w_norm: 1.8420240652914215	residual_error: 92.39630102542934	lambda:
16.635533333438687		
w_norm: 1.7578344358851319	residual_error: 92.2403313247562	lambda:
17.32868051399863		
w_norm: 1.6912694674229027	residual_error: 92.11599146101157	lambda:
18.021827694558574		
w_norm: 1.5822541740103822	residual_error: 91.92717363184512	lambda:
18.714974875118518		
w_norm: 1.426055638400412	residual_error: 91.72082256986741	lambda:
20		



b)

```
[246]: error_rate = []
        sparsity = np.zeros(len(lambda_array))
        for i in range(len(lambda_array)):
            diff = abs(np.sign(X_train@w[:,i]) - y_train[:,0])
            for j in w[:,i]:
                # Add to count the number of nonzero entries of w => if |w|_i > 1e-6
                if j > 1e-6:
                    sparsity[i] += 1
            error_rate.append(diff.sum()/len(diff))
        plt.scatter(x=sparsity, y=error_rate)
        plt.xlabel("Sparsity")
        plt.ylabel("Error rate")
        plt.show()
        # As the number of nonzero entries in w increases => the error rate decreases
```



c)


```
[247]: X_test = X[101:]
       y_test = y[101:]
```

```
[253]: # a)
       residual_error = []
       w_norm = []
       for i in range(len(lambda_array)):
           residual_error.append(np.linalg.norm(X_test@w[:,i] - y_test))
           w_norm.append(abs(w[:,i]).sum())
           print("w_norm: " , w_norm[i], "\tresidual_error: ", residual_error[i],
               ↳ "\tlambda: ", lambda_array[i])
       plt.scatter(x=w_norm, y=residual_error)
       plt.show()
       # As with the train data, as lambda increases the residual error and the norm-1
       ↳ of w decreases.
       # However, the plot show a more curved graph with it's point with a more
       ↳ distance between them
```

w_norm: 36.109055952630406	residual_error: 201.21304378897253	lambda: 1e-06
w_norm: 18.1266980593891	residual_error: 197.91570675411856	lambda: 0.6931481805599453
w_norm: 13.51150995902426	residual_error: 195.93637747867436	lambda: 1.3862953611198905

w_norm: 11.207473857087024 2.079442541679836	residual_error: 194.34688419042368	lambda:
w_norm: 9.600042255143801 2.7725897222397813	residual_error: 192.82475715297167	lambda:
w_norm: 8.425137987735571 3.4657369027997267	residual_error: 191.78903343197146	lambda:
w_norm: 7.504537201776889 4.158884083359672	residual_error: 190.782753642679	lambda:
w_norm: 6.785326085022339 4.852031263919617	residual_error: 189.77979622919315	lambda:
w_norm: 6.123578969469852 5.545178444479562	residual_error: 188.87130033785922	lambda:
w_norm: 5.579710963832548 6.238325625039508	residual_error: 187.91498643542022	lambda:
w_norm: 5.137386633432994 6.931472805599453	residual_error: 187.08197273897952	lambda:
w_norm: 4.7294810562842615 7.624619986159399	residual_error: 186.45962974159804	lambda:
w_norm: 4.355999922114707 8.317767166719344	residual_error: 185.78010100452067	lambda:
w_norm: 3.9723493979938773 9.01091434727929	residual_error: 185.0836551926895	lambda:
w_norm: 3.75616629710407 9.704061527839235	residual_error: 184.57120572253032	lambda:
w_norm: 3.4778926816773117 10.39720870839918	residual_error: 183.74603478027777	lambda:
w_norm: 3.136669482594097 11.090355888959126	residual_error: 182.7934822301498	lambda:
w_norm: 2.9903158145972366 11.783503069519071	residual_error: 182.30658498752715	lambda:
w_norm: 2.764883558007083 12.476650250079016	residual_error: 181.60219752866593	lambda:
w_norm: 2.5863002243014237 13.169797430638962	residual_error: 181.06824398874377	lambda:
w_norm: 2.4504933690532065 13.862944611198907	residual_error: 180.7230201718312	lambda:
w_norm: 2.2633392534203525 14.556091791758853	residual_error: 180.17953711686815	lambda:
w_norm: 2.0540384197405395 15.249238972318798	residual_error: 179.62115975839671	lambda:
w_norm: 1.956318032917396 15.942386152878743	residual_error: 179.41276239881375	lambda:
w_norm: 1.8420240652914215 16.63553333438687	residual_error: 179.11612841938015	lambda:
w_norm: 1.7578344358851319 17.32868051399863	residual_error: 178.87761119660146	lambda:
w_norm: 1.6912694674229027 18.021827694558574	residual_error: 178.69479609154843	lambda:

w_norm: 1.5822541740103822 residual_error: 178.4147295472156 lambda:
 18.714974875118518
 w_norm: 1.426055638400412 residual_error: 178.11348094047474 lambda:
 20



```

[252]: # b)
error_rate = []
sparsity = np.zeros(len(lambda_array))
for i in range(len(lambda_array)):
    diff = abs(np.sign(X_test@w[:,i]) - y_test[:,0])
    for j in w[:,i]:
        # Add to count the number of nonzero entries of w => if |w|_i > 1e-6
        if j > 1e-6:
            sparsity[i] += 1
    error_rate.append(diff.sum()/len(diff))
plt.scatter(x=sparsity, y=error_rate)
plt.xlabel("Sparsity")
plt.ylabel("Error rate")
plt.show()
# Now, the error rate vs sparsity behaves completely different as before. As
# the sparsity increases, the error
# rate increases with it.
  
```

